ACTIVELY AI

Deep-Dive Company Research

April 2026 · Prepared for internal use

CONTENTS

01 — Company at a Glance
02 — Founding Story & Origin
03 — Product: How It Actually Works
04 — Business Model & Positioning
05 — Funding & Investors
06 — Customer Evidence
07 — Competitive Landscape
08 — Team & Culture Signals
09 — Open Questions & Risks
10 — Growth Initiative Ideas [annotate here]

01 — Company at a Glance

$22.5M	2022	~50
Total Raised	Founded	Employees (est.)

Actively AI is a New York-based startup building what it calls "GTM Superintelligence" — AI that doesn't just automate sales tasks, but reasons through them. The company sits at the intersection of two major market forces: the explosion of AI reasoning models (think OpenAI o1, DeepSeek) and the collapse of the old volume-based outbound playbook that powered a decade of B2B sales.

In plain terms: where most AI sales tools help reps send more emails faster, Actively helps reps figure out which emails should ever be sent at all — and why, and when. The company serves enterprise and high-growth companies, with marquee customers including Ramp, Ironclad, Greenhouse, and Verkada.

FAST FACTS

- **Founded:** Founded:
  - 2022

Headquarters	New York, NY
Founders	Mihir Garimella (CEO) & Anshul Gupta (Co-Founder)
Investors	Bain Capital Ventures (lead), First Round Capital, Lachy Groom, Stanford AI faculty
Customers	Ramp, Ironclad, Greenhouse, Verkada, Justworks, Attentive
Category	GTM AI / Revenue Intelligence / AI Sales Agent

02 — Founding Story & Origin

Who Are the Founders?

Actively was co-founded by Mihir Garimella (CEO) and Anshul Gupta, both Stanford AI researchers. Importantly, neither came from a sales background — they came from math, computer science, and AI research, and have been studying and building in the AI space since GPT-1, long before the current generative AI wave.

MIHIR GARIMELLA — CEO

Mihir's Stanford research focused on active learning — a branch of machine learning where models identify which data points are most valuable to learn from, rather than training on everything indiscriminately. This is literally the intellectual origin of Actively: the idea that intelligent selection beats brute-force processing. Active learning is also where the company's name comes from.

Mihir has been outspoken in his belief that AI reasoning — not AI volume — is the next frontier of GTM. He's published thought leadership on human-agent collaboration for long-running AI agents, a topic that maps directly to Actively's architecture.

ANSHUL GUPTA — CO-FOUNDER

Anshul holds an Integrated Dual B.Tech/M.Tech from IIT Roorkee, with deep experience spanning Microsoft Azure (Cloud + AI, Identity & Cybersecurity), the Federal Reserve (ML for macroeconomic forecasting), and a Presidential executive order on AI — a remarkably broad background for a startup founder. Before Actively, he worked in financial modeling and investments, giving him an unusual combination of technical depth and commercial instinct.

THE ORIGIN INSIGHT

Both founders observed the same problem from the research side: B2B selling had become a high-cardinality, high-dimensionality problem — meaning the number of variables a rep needs to process to make a good decision (which product? which persona? which signal? which timing?) had exploded far beyond human capacity. But the software tools hadn't caught up. Most tools either drowned reps in data signals or auto-sent generic outreach at scale.

> *"What if you could assign 1 AE to each account in your TAM? The depth with which they'd look to research, understand, and constantly figure out how to best engage that account is what we aim to replicate with AI."*
> **— Anshul Gupta, Co-Founder**

The company went through Y Combinator, where the founders say the experience fundamentally changed how they thought about culture and speed — treating culture not as perks and values on a wall, but as "how fast your team can move without breaking."

WHY THEY NAMED IT "ACTIVELY"

The name has a double meaning: active learning (the AI technique from Mihir's research) and the idea that the system is always-on, always learning, always getting better. As they put it: "Actively

stands for active learning — a concept from the founders' research background at Stanford. Most GTM tech gets worse the more you use it. Actively gets better."

03 — Product: How It Actually Works

Actively's core product is a per-account AI agent system that runs 24/7 across every account in a company's TAM (Total Addressable Market). Each morning, the system analyzes each account and surfaces prioritized recommendations with explicit reasoning — not just a ranked list, but a hypothesis for why to engage now, how to approach, and what to say.

THE FIVE-LAYER ARCHITECTURE

Layer 1: Custom Model Training

Unlike generic AI tools, Actively builds a private, custom reasoning model for each customer. This model is trained on the company's own data: CRM records, closed-won patterns, Gong call recordings (including specific moments — they cite analyzing "the 27th minute of a Gong call"), Salesforce custom objects, product usage data, and more. The model learns what actually converts for that specific company, not just generally.

Layer 2: Signal Synthesis

Actively pulls from first-party data (CRM, call transcripts), second-party data (intent signals from tools like 6sense, ZoomInfo), and third-party data (job postings, news, funding announcements, public filings). It synthesizes all of this into a single "context window" per account — a complete picture that no individual rep could assemble manually.

Layer 3: Reasoning & Hypothesis Generation

This is the core differentiator. Rather than producing a score, Actively's reasoning agents form explicit hypotheses: "This account is worth engaging now because their VP of Sales posted about pipeline efficiency challenges, they just hired 3 new AEs, and their Gong transcripts show repeated objections around rep ramp time." The system uses reasoning models from OpenAI and Anthropic, combined with in-house models, to generate these hypotheses.

Layer 4: Action Delivery

Recommendations are delivered inside the tools reps already use (Salesforce, Outreach, etc.) with pre-drafted messaging, contact-level strategies for multi-threading, and call scripts. Reps don't need to leave their workflow. Each recommendation comes with a "why now" explanation so reps understand the reasoning, not just the output.

Layer 5: Active Learning Loop

Every interaction — replies, meetings booked, deals won or lost — feeds back into the model. The system continuously refines its hypotheses and improves its predictions. This creates a compounding advantage: the longer a customer uses Actively, the better the model gets at predicting what works for that specific company.

KEY USE CASES

- **Daily outbound workflow:** Daily outbound workflow:
  - Prioritized account/contact lists with "why now" reasoning delivered each morning

- **Territory planning:** Territory planning:
  - AI-powered account distribution based on win probability, not just geography
- **Competitor takeouts:** Competitor takeouts:
  - Campaigns targeting accounts currently using competitor products
- **Cross-sell/upsell:** Cross-sell/upsell:
  - Identifying expansion opportunities inside the existing customer base
- **New rep onboarding:** New rep onboarding:
  - New hires contribute pipeline from day 8 by working Actively's recommendations immediately

THE GOOGLE PAGERANK ANALOGY

> *"Yahoo and MSN had search, but they treated all keywords equally. Then Google introduced PageRank — a system that mined through the web's link structure to find what actually mattered to users. It didn't win with a better UI. It won by solving the right problem. That's what Actively is doing for GTM."*
> **— Anshul Gupta, Co-Founder**

04 — Business Model & Positioning

Target Customer

Actively sells to enterprise and high-growth B2B companies with significant sales development functions — companies with enough outbound volume that the "quality vs. quantity" problem is acute. Key buyer personas are CROs, VP Sales, VP of Sales Development, and Revenue Operations leaders. Their sweet spot appears to be companies in the $50M–$1B ARR range, though they have enterprise customers as well.

Pricing Model

Actively is not public with pricing, but based on their customer profile and the depth of customization involved (per-company model training, CRM integration, ongoing active learning), this is almost certainly a SaaS subscription with enterprise-level ACV — likely $50K–$300K+ annually per customer, depending on team size and scope.

Go-To-Market Motion

Actively sells top-down, targeting C-suite and VP-level buyers with a strong proof-of-value motion. They offer a proof-of-concept period (Ironclad noted that "reps booked quality meetings on day one" during the POC). Customer evidence — case studies, testimonials, and quantified ROI — is central to their sales process.

THEIR POSITIONING NARRATIVE

Actively positions against the entire category of "AI SDR" and "AI sales assistant" tools, arguing that those products failed because they optimized for volume. Their core narrative:

- "More is not more." The volume playbook is broken — inboxes are flooded, response rates are collapsing.
- The best sellers don't win by doing more. They win by reasoning better — deeper research, precise hypotheses, perfect timing.
- Actively encodes that reasoning into AI, making every rep perform like your top performer.
- This is not a copilot. It's superintelligence — it makes decisions, not suggestions.

This narrative is aggressive and deliberately provocative. It doesn't just differentiate — it indicts the entire competitive category. That's a high-risk, high-reward positioning strategy that works when you have the customer proof to back it up. They appear to have that proof.

Defensibility

Actively's moat deepens over time in two ways: (1) the per-company custom model improves with each interaction, creating switching costs as customers accumulate trained intelligence; and (2) the institutional knowledge problem they solve becomes more acute at larger companies, making churn

increasingly painful. This is a compounding data advantage, which is the strongest kind of moat in enterprise AI.

05 — Funding & Investors

$5M	$17.5M	$22.5M
Seed — First Round Capital	Series A — Bain Capital Ventures	Total Raised

Actively announced their Series A on April 2, 2025. The round was led by Bain Capital Ventures, with continued participation from First Round Capital (who led the seed). The seed round was previously unannounced, making the Series A announcement a dual reveal.

WHY THESE INVESTORS MATTER

Bain Capital Ventures

BCV is not just a financial backer — they are strategic validators. Their GTM tech portfolio includes Apollo (one of the largest sales intelligence platforms), 6sense (intent data leader), and Clari (revenue forecasting). BCV evaluated, by their own account, "every AI sales technology" before betting on Actively. This is a strong signal that Actively is the category winner in their view, not just a promising entrant.

First Round Capital

First Round is one of the most prestigious seed-stage firms in tech (seed investors in Uber, Square, Roblox, Clearbit). Their continued participation in the Series A is a positive signal — seed investors who double down at A are making a bet that the company has achieved meaningful validation. Notably, Clearbit — a data enrichment company for B2B sales — was also a First Round company, suggesting familiarity with the GTM data space.

Lachy Groom

Lachy Groom is a former Stripe executive turned angel investor, known for backing fintech and enterprise software companies. His involvement suggests Actively may be targeting fintech-adjacent companies (consistent with their Ramp customer relationship) and brings network value in those verticals.

Stanford AI Faculty

Having Stanford AI faculty as investors is unusual and meaningful — it signals academic credibility, provides access to top AI research talent for recruiting, and suggests the underlying technical approach is sound from a research perspective.

06 — Customer Evidence

Actively's customer case studies are unusually specific — they include hard revenue numbers and percentage lifts, not just qualitative testimonials. This is notable because it suggests strong customer willingness to share results publicly, which usually indicates high satisfaction and low churn risk.

Ramp

Ramp is Actively's marquee reference customer. In 12 months, Ramp attributed tens of millions in net-new closed-won revenue to the Actively system. Win rates for Actively-sourced opportunities are 22% higher than traditionally-sourced opportunities. Every new rep at Ramp is now onboarded on Actively from day one — meaning it's embedded as core infrastructure, not an add-on tool. As Ramp continues to rapidly launch new products, Actively's intelligence engine helps them identify the most relevant buyers for each product.

> *"We have multiple data and signal feeds at Ramp. Actively AI not only activates them for our Sales team but presents them with complete context in actionable, timely plays that convert to quality pipeline."*
> **— Sales Leader, Ramp**

Greenhouse

Greenhouse (HR tech, 800+ employees, SMB/MM/Enterprise segments) deployed Actively across their global SDR team — US, EMEA, and APAC. Results were tracked against the 6 months before implementation vs. 6 months after.

+54% Pipeline per rep	**+58%** Conversion rate	**50%** Reduction in ramp time

They also generated $12.8M in directly attributed qualified pipeline. Onboarding time was cut from 3.5 weeks to 2 weeks, and 100% of new hires now create opportunities during their onboarding month.

Ironclad

Ironclad (legal tech, CLM platform serving L'Oréal, OpenAI, Cisco) used Actively to address a specific challenge: flat headcount during a growth period. Their BDR team had to cover more accounts without adding people. Actively allowed them to convert contacts to qualified pipeline at a 41% higher rate, and within six months they attributed tens of millions in qualified pipeline to the

system. Ironclad now uses Actively for both inbound qualification and competitive takeout campaigns (targeting companies using competing CLM tools).

Verkada

Verkada (enterprise security tech) uses Actively for their Managed Detection & Response (MDR) program. Their Sr. Director called it "the future of outbound" and specifically noted that it would be "a model for every ambitious growth stage company." This testimonial is useful for Actively's own sales process — Verkada's success story maps well to the mid-market and security verticals.

07 — Competitive Landscape

The AI sales technology market is crowded, but Actively's positioning is designed to sidestep most competitors by changing the frame entirely — from "AI sales tool" to "GTM superintelligence." Here's how the main competitors stack up:

Clay

APPROACH

Data enrichment and waterfall lookup tool. Clay lets users build workflows that pull from many data sources to enrich leads. Highly flexible, beloved by RevOps power users.

KEY WEAKNESS VS. ACTIVELY

Clay is a tool, not intelligence. It surfaces data; it doesn't reason about it. Clay tells you that a prospect hired 5 engineers — Actively tells you why that matters for your specific product and when to reach out. Clay requires significant manual setup and expertise to use well, creating an ops overhead problem.

ACTIVELY'S EDGE

Actively has the reasoning layer Clay lacks. Where Clay requires a RevOps engineer to build and maintain workflows, Actively's custom model handles the reasoning automatically and continuously improves.

Artisan / 11x / AI SDR tools

APPROACH

Fully automated AI SDRs that send outbound emails and make calls at high volume. These tools promise to replace SDRs entirely with AI agents.

KEY WEAKNESS VS. ACTIVELY

These tools are the category Actively explicitly indicts. They optimize for volume, which floods inboxes with generic outreach and tanks reply rates. As Anshul Gupta told TechCrunch, AI SDRs "failed" by focusing on volume. Actively's approach is the opposite: fewer, higher-quality touchpoints backed by actual reasoning.

ACTIVELY'S EDGE

Actively is not trying to replace reps — it's trying to make reps superhuman. This positioning is more palatable to buyers (less threatening) and produces better outcomes (human relationship + AI intelligence).

Apollo.io

APPROACH

Sales intelligence and engagement platform with a large contact database, sequences, and basic AI features. Very popular at the SMB/mid-market level.

KEY WEAKNESS VS. ACTIVELY

Apollo is a database + sequencing tool, not a reasoning engine. Their AI features are largely generic — email personalization, basic contact recommendations. They're built for volume and breadth, not depth.

ACTIVELY'S EDGE

Actively's per-company custom model is fundamentally different from Apollo's generic AI. Actively also partners with and sits on top of data tools like Apollo — they're not direct substitutes at the enterprise level.

Gong / Chorus (Revenue Intelligence)

APPROACH

Conversation intelligence tools that record and analyze sales calls. Gong provides insights into pipeline health and rep performance based on call data.

KEY WEAKNESS VS. ACTIVELY

Gong is retrospective — it tells you what happened. Actively is prospective — it tells you what to do next. Gong also doesn't handle prospecting or outbound prioritization.

ACTIVELY'S EDGE

Actively actually ingests Gong data as an input to its model. They're complementary in the stack, not competitive. Actively uses what Gong records to improve its recommendations.

Salesforce Einstein / HubSpot AI

APPROACH

CRM-native AI features built into existing platforms. Salesforce Einstein offers lead scoring, predictive forecasting, and basic recommendations.

KEY WEAKNESS VS. ACTIVELY

CRM-native AI is limited by CRM data — it only knows what's in Salesforce. Actively synthesizes CRM data plus external signals, call recordings, intent data, and more into a richer picture. CRM AI also isn't reasoning-based; it's statistical scoring.

ACTIVELY'S EDGE

Actively can sit on top of any CRM and use CRM data as an input. They're not competing for the CRM itself — they're adding a reasoning layer on top.

08 — Team & Culture Signals

Understanding the team DNA matters if you're thinking about working there. Several signals stand out from public information:

Technical Depth + Commercial Instinct

The founders' combination of Stanford AI research backgrounds with real commercial experience (Federal Reserve, Microsoft, financial modeling) is rare. BCV's investor note specifically called this out: "It's rare to meet someone who can explain the nuances of RAG, while also knowing the system behind MEDDPICC." This combination tends to produce founders who can both build and sell — a strong predictor of durable company success.

Culture Language

Actively's careers page articulates a set of values worth reading carefully:

- "Seek truth and act with rigor and logic. Exceptional products aren't built on assumptions."
- "Time is a scarce resource. We act with urgency, hunger, and precision."
- "Own outcomes — do whatever it takes to see an initiative through, including the last mile."
- "Think outside the box and look for exponential improvement. Don't accept incrementally."
- "We deeply understand our customers' business goals and relentlessly focus on driving outcomes."

This is not boilerplate startup culture language. The emphasis on rigor, urgency, and customer outcomes suggests a company that operates with high standards and moves fast. The "including the last mile" phrasing is notable — it implies they have low tolerance for near-misses or hand-offs that don't complete.

Blog as Signal

Actively's blog is technically substantive. Recent posts from the Head of Engineering include "Designing Multi-Agent Systems at Enterprise Scale" and "Future of AI Agents: Lessons from Actively, Decagon, and Hebbia." Co-founder Anshul wrote "Building Context Graphs for GTM — and Why Salesforce Can't Do It." These are not marketing fluff — they're technical artifacts that reflect actual thinking. Companies that publish this kind of content tend to have strong intellectual culture internally.

Team Composition

The team includes AI researchers and engineers from Google, Scale AI, Ironclad, Waymo, Robust Intelligence, and Microsoft, alongside GTM leaders who have scaled orgs from early stage to $100M+ ARR. This mix of deep technical talent and practical GTM operators is unusual — most AI sales startups are either too technical (no GTM instinct) or too sales-focused (AI is shallow). Actively appears to have both.

09 — Open Questions & Risks

No company research is complete without an honest look at what's uncertain or potentially risky. These are the questions worth thinking about:

1. Enterprise Sales Cycle Length

Custom model training, deep CRM integration, and proof-of-concept periods suggest a long sales cycle — potentially 3–6 months from first contact to close. For a company at Series A stage, this creates cash flow pressure and limits how quickly they can grow customer count. The question is whether they can streamline onboarding without sacrificing the customization that makes the product work.

2. The "Custom Model" Scalability Question

Building a bespoke reasoning model for each customer is expensive and time-intensive. As they scale to hundreds of customers, how do they maintain quality? Do they template the process? Use fine-tuning rather than full custom training? The scalability of their core value proposition is an open technical and operational question.

3. Competition from Foundation Model Companies

OpenAI, Anthropic, and Google are all building increasingly capable reasoning models. As these models improve, the bar for what "reasoning" means will rise. Actively's moat isn't just in the reasoning — it's in the per-company data and the active learning loop. But if foundation models commoditize enough of the reasoning layer, Actively may need to rely more heavily on their data network effects.

4. SMB vs. Enterprise Focus

All of Actively's public customer logos are large or high-growth companies (Ramp, Ironclad, Greenhouse, Verkada). This is a natural starting point, but it raises questions about whether their product works for smaller companies with less historical data to train on. Do they plan to move downmarket? Or does their business model require enterprise customers indefinitely?

5. Buyer Consolidation Risk

Salesforce, HubSpot, and Microsoft are all building AI features natively into CRMs. If any of them can credibly deliver reasoning-based prioritization inside their existing platforms, Actively's standalone value proposition gets harder to sell. Their best defense is continued performance differentiation and deep data integrations that CRM-native tools can't match.

10 — Growth Initiative Ideas

[Use this page to annotate, react, and build on these ideas]

01 Vertical-Specific Proof Packages

THE IDEA

Actively's case studies span fintech (Ramp), legaltech (Ironclad), HR tech (Greenhouse), and security (Verkada) — but they're packaged generically. Build lean, vertical-specific one-pagers that translate Actively's metrics into the language of each target buyer. A VP of Sales at a devtools company doesn't see themselves in the Ramp story. Give them a story that sounds like their world.

WHY NOW

Actively just closed their Series A and is scaling fast. They need to expand into new verticals efficiently. Vertical proof packages let the sales team reach new buyers without rebuilding the sales narrative from scratch each time.

✏ *Which verticals do you think are highest priority? What data would you want in a vertical one-pager?*

---

02 Competitive Displacement Campaign Engine

THE IDEA

Actively's narrative is that AI SDRs "failed" — yet they're not systematically going after accounts currently paying for those tools. Build a displacement playbook: identify ICP accounts using Artisan, 11x, or Clay for outbound, and run targeted campaigns that make the "volume vs. reasoning" contrast visceral and personal. Use Actively's own technology to power the campaign — the meta-story sells itself.

WHY NOW

Post Series A, Actively needs to grow ARR aggressively. Displacement campaigns are the fastest path to revenue when you have a credible performance narrative. The AI SDR category is also under scrutiny — timing is right for buyers already having doubts.

✏ *Is this already happening? What does the current competitive campaign strategy look like?*

---

03 "New Rep Accelerator" as a Category Narrative

THE IDEA

Greenhouse's results showed 50% faster onboarding and 100% of new hires creating pipeline in their first month. Ramp trains every new rep on Actively from day one. This is actually a distinct product use case — Actively as an onboarding infrastructure tool, not just a prospecting tool. Build a dedicated "New Rep

Accelerator" narrative targeting companies experiencing high rep turnover (average rep tenure is 1.5 years per Actively's own data) or rapidly scaling their sales teams.

WHY NOW

Rep turnover is a top pain point for CROs. Companies that are hiring aggressively (post-funding, post-product launch, entering new markets) are acutely aware of how long it takes new reps to ramp. This narrative opens a new buying trigger that doesn't require existing pipeline pain.

✏ *Is there a hiring-signal trigger you're already using to find this buyer? What's the current ICP qualification criteria?*

---

04 Customer Expansion Playbook for Existing Logos

THE IDEA

Actively's per-account agents work on cross-sell and upsell too — but this use case feels underleveraged in their public messaging. Build a dedicated expansion playbook for existing customers: starting with SDRs, then expanding to AEs (for deal advancement), then AMs (for upsell/cross-sell), then RevOps (for territory planning). Each expansion is a land-and-expand motion within the existing account.

WHY NOW

Their current customer base (Ramp, Greenhouse, Ironclad) are large companies with many teams. Getting a second or third team deployed is significantly cheaper than acquiring a new customer and contributes to ACV expansion and NRR.

✏ *What does the current expansion motion look like? Is there a dedicated CSM/growth function or does sales own expansion?*

---

05 Actively's Own Outbound as a Living Demo

THE IDEA

Anshul Gupta mentioned in a podcast that Actively uses poor cold outbound campaigns directed at them as a trigger signal for their own outreach — essentially using bad AI SDR emails they receive as intent data. Lean into this more aggressively and publicly. The meta-narrative of "we use our product to sell our product, and here's the result" is extraordinarily credible and makes for compelling content. Run a transparent campaign series: show how Actively identifies prospects, what hypotheses it generates, and what results it drives — for Actively's own growth.

WHY NOW

Buyers are skeptical of AI tools that don't eat their own cooking. A public, transparent account of Actively using Actively to grow Actively would be the most credible proof point they could produce — and it generates content, brand awareness, and sales conversations simultaneously.

✏ *Is there appetite for this level of transparency? What would the legal/competitive risk be of sharing this publicly?*
